


Data Engineering Assessment - Matheus Alves



The Challenge

1. By making use of the **invoices.xls** data file as your source, create a Data Warehouse model with best practices in place, using the Kimball methodology.
2. Populate the model created in step 1 with the data provided in the Excel sheet.
3. Analyze any abnormalities (if any) in the data provided and take any action needed (where possible).
4. List any assumptions taken into consideration.
5. Provide 3 different aggregations one might use for reporting purposes.



Table of Contents

- 0. Base Constraints
- 0.1. Scope Matrix - Documentation and Code
 - 0.1.1 Base General Process
 - 0.1.2 Base ETL Considerations
 - 0.1.3 Base Warehouse Considerations
- 1. Assumptions & Abnormalities
 - 1.1 Base Considerations
 - 1.2 Overall Data Assumptions
 - 1.3 Overall Data Abnormalities
 - 1.3.1 Country column
 - 1.3.2 Customer ID column
 - 1.3.3 Price column
 - 1.3.4 Quantity column
 - 1.3.5 Description column
 - 1.3.6 StockCode column
 - 1.3.7 General columns data
- 2. Pipeline Architecture Diagram
- 2.1 Data Warehouse Schema Diagram
- 2.2 Documentation
 - Tables and Columns
- 3. Packages and modules overview
 - 3.0. Running the project

- 3.1. Main Script
 - 3.2. infra
 - 3.3. ingestion
 - 3.4. utils
 - 3.5. assets
- 4. General Code Structure
- Aggregations & Reporting
 - 1. Insight - (AVTQ) Absolute Value of Transaction Quantities and Revenues across different categories
 - 2. Insight - (SLICR) Sales Location Impact Comparison on Revenue
 - 3. Insight - (CLV) Customer Lifetime Value
 - 4. Data-Driven Deep
- Tunning Scenario

0. Base Constraints

- This task may be done either with MSSQL, Python or C#.
- It is highly recommended that comments are included, and code is structured well.
- If making use of python, you can make use of Jupyter notebooks, thus you can either
- provide a .PY file or .IPYNB file.
- Assumptions & Abnormalities are to be provided in a separate document.

0.1 Scope Matrix - Documentation and Code

Base General Process

Feature	In Scope (MVP)	Out of Scope
README.md-based documentation (iterative view)	✓	
PDF documentation (static)	✓	
Module and package details	✓	
Base real case assumptions	✓	
Packages and modules partially/fully reusable	✓	
Base data quality check	✓	
Cloud-related environments		✓
Full production-ready		✓
Non-Python frameworks for ETL		✓
Specific metrics logging (to be used in Prometheus, for instance)		✓
Deeply statistic approaches in warehouse/data cleaning		✓

Feature	In Scope (MVP)	Out of Scope
Specific Architecture Consumption/Tunning Scenario		✓

Base ETL Considerations

Feature	In Scope (MVP)	Out of Scope
Process diagrams	✓	
Isolated environment (.venv, optional)	✓	
pip requirements (pip install -r requirements_xx.txt)	✓	
APIs		✓
Scheduler-ready		✓

Base Warehouse Considerations

Feature	In Scope (MVP)	Out of Scope
Data full metadata	✓	
Client basic tuning example (MSSQL Query Plan)	✓	
Decision explained	✓	
Physical computing/database architecture and setup		✓
Server basic tuning (on possible queries)		✓

1. Assumptions & Abnormalities

1.1 Base Considerations

There are two main logic on this assesment answers.

- **draw.ipynb (Jupyter Notebook):** Containing draws.
- **solution.py (Script):** Pipeline solution. Consideration:
 - Assumptions usually are discussed with the areas/teams involved. In a real-world scenario, I would discuss these assumptions with the business team to ensure that they are aligned with the business rules and requirements.
 - **This consideration applies to every main assumption I make.**
- This code is made on python + mssql.
- New data overwrite old data.

1.2 Overall Data Assumptions

- I'm assuming that the data is a snapshot and we don't need to update it in the database.
- Data snapshot only; no updates required.
 - Even with that in mind, we have dimensions, for instance, that take into consideration Slowly Changing Dimensions (SCD) Type 2, to keep track of the changes in the data.
- Lowercase will be used for column names, in addition to snake case.
 - This is a common practice to avoid any possible issues with the database.
- Adjust data.
 - Maintained and flagged, proper differentiated from sales.
 - adjustments can address wastes in Lean Methodologies as well potential losses by possible mistakes, and can come to be explored.
- I am assuming that every valid product adjustment at any category is being properly identificate in the description column.
- I am assuming that every known retail is registered, and related transactions share the same ID across different lines.
- I am assuming that I am dealing with hybrid store sales. That sails are made both online and offline.
- A simple solution to mapped values is to use a dictionary to standardize the values.
 - More complex cases involving large datasets, approaches can range from using a Cartesian mapping stored in a .txt file to developing a classification model to identify and standardize variations in location names automatically.
- No special cases in invoice date column.
- Warehouse inserted data was also approached with null scenarios. They were ignored, as they was threatened as special cases. I'm assume I've maintained referencial integrity.

1.3 Overall Data Abnormalities

- Columns contain missing values that can represent various scenarios, ranging from errors/bugs to information that was not shared.
 - I am categorizing these values appropriately.
 - Unspecified: Explicitly chose not to declare their information.
 - Null (e.g., None, NaN): Did not provide the information at all.
- I am assuming this governance rule (when I have not specified a different approach):
 - Placeholder values will be standardized as "Unspecified."
 - Null values will remain as null.
- There were test values that have been addressed and removed from all columns.
 - I am assuming that there are no special cases, such as identifying errors in the client flow that need to be registered.

- To effectively address description and category problems, I would evaluate the possibility of creating a governance policy to map such cases according to the business rules.
 - I'm assuming that this scenario involves areas where I am unable to act.

1.3.1 Country column

- Some entries in the Country column might not represent valid country names (e.g., typos, placeholders, or regions like "Channel Islands").
 - The simplest approach would rename Country to Locations.
- Some entries in the Country column are abbreviated, such as "United Kingdom" and "UK".

1.3.2 Customer ID column

- Test customers identified, when casually understanding data format.
 - There is 'test' data on different columns.
- Null clients often relates (based on description column) to manual inserts, fee, a/b testing. I'll deal with them in the following lines.
- Sales that doesn't have a description, as well doesn't have a customer ID, will be flagged as special cases.
- I'm assuming that customer id and description missing, being < 0, are error entries.

1.3.3 Price column

- Prices equals to zero or negatives.
 - There were unusual descriptions such as "mixed" and "short". They are ignored, as they are relevant sales.
 - I will map returns and consider them valid only when Quantity_return is less than Quantity_sale (based on merging invoice and StockCode data, validating negative values against all related entries).
- Values involving different scenarios.
 - Damages, damaged, bad quality.
 - Damages or negative descriptions appointments at any kind will be flagged as lost sales.
 - Descriptions that are not clear or relevant, such as 'lost in space,' will be flagged as lost sales. I will assume these represent some kind of product damage or loss.
 - I am also assuming that there are no valid descriptions that share the same part of their definition. For instance, terms like "wet" in descriptions such as 'WET PLANT PERFECTLY DECORED' are not being used ambiguously.
 - I'm using a list of known descriptions patterns to flag lost sales.

1.3.4 Quantity column

- Negative and zero values that isn't some returned data.
 - Some negative data refers to discounts. They are flagged as discount.
- There where match in stockcode agains't null values. For rows without a Description or Price, I am assuming the last recorded Description and Price based on the StockCode. This assumes that StockCode descriptions are static, and the last recorded value is correct and applicable

1.3.5 Description column

- There are ids on description, that seems to be a stockcode, but there aren't any indicator. I'm assuming that they are
- AMAZON FEE, adjustments in bad debt, etc, they will be flagged as financial detail.
- Adjustments are flagged as maintenance data.
 - Only adjustments proper described will be flagged as maintenance data.
- Update are flagged as maintenance data.
- Bargains are discounts, but they are threat directly as financial detail.

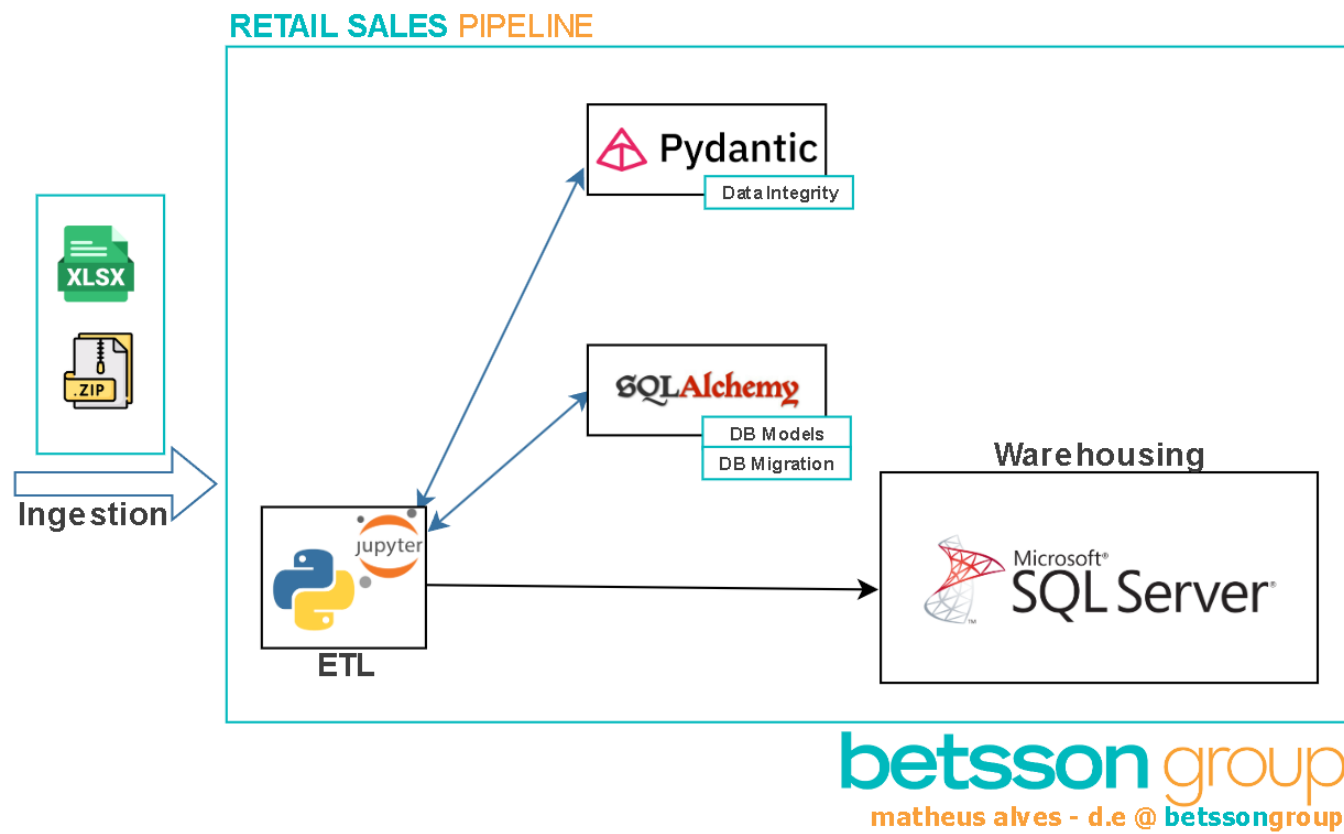
1.3.6 StockCode column

- There is patterns in stockcode like "M", "D". Coming exclusively, they are representating discount. Manual sales are sales as the others, they are ignored.
- Discounts are flagged as discount.
- Stock code S, representing samples, typically has values of -1. However, there are some cases where this is not the case. These will be flagged to be removed.
- Gift are flagged as lost sales
- Test are flagged to be removed.
- Stockcode flagged on DW when. without return flag or quantity < 0.
- Bankcharges are flagged as financial detail.

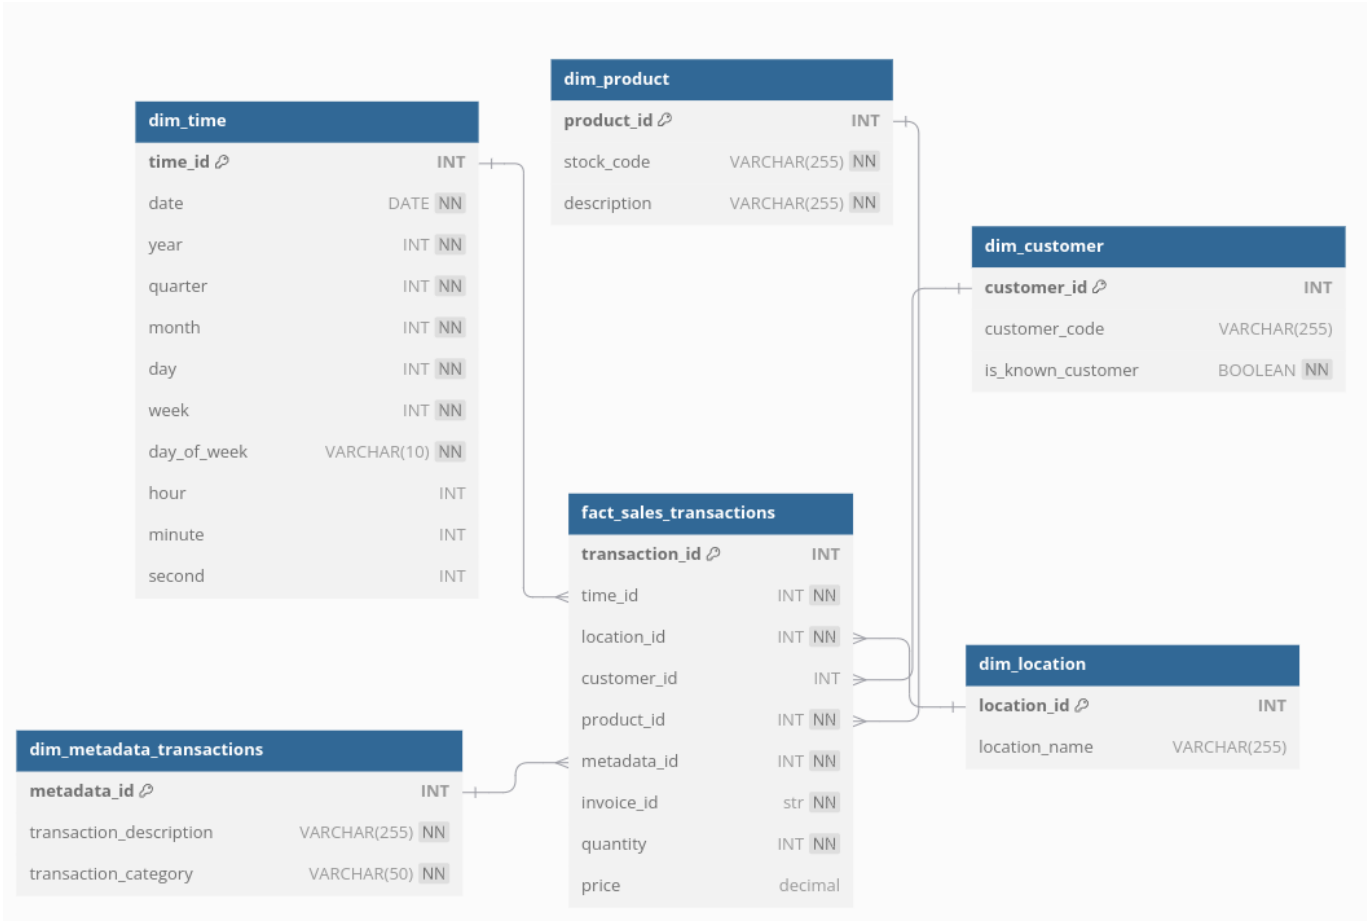
1.3.7 General columns data

- Guarantee that overall mapped appointment is being properly flagged.
- After filtering all flagged columns, any remaining abnormal data is reviewed.
- I've tried to enrich data, but there were no match:
 - Price: Updated using the most recent value for the same stock_code.
 - Description: Updated using the modal (most frequent) value for the same stock_code.
 - Assuming that if resulting value is null and there is no client id, so it is entry error.

2. Pipeline Architecture Diagram



2.1 Data Warehouse Schema Diagram



2.2 Documentation

- Anonimized data.

Tables and Columns

dim_time

Column	Description
time_id	Primary key for the time dimension.
date	The full date in YYYY-MM-DD format.
year	The year of the transaction.
quarter	The quarter of the year (1-4).
month	The month of the year (1-12).
day	The day of the month (1-31).
week	The week of the year (1-53).
day_of_week	The name of the day (e.g., "Monday").
hour	The hour of the transaction (0-23, optional).
minute	The minute of the transaction (0-59, optional).
second	The second of the transaction (0-59, optional).

dim_location

Column	Description
location_id	Primary key for the location dimension.
location_name	The name of the location where the transaction occurred.

dim_customer

Column	Description
customer_id	Primary key for the customer dimension.
customer_code	Unique code identifying the customer, if available.
is_known_customer	Indicates whether the customer is known (True) or anonymous (False).

dim_product

Column	Description
product_id	Primary key for the product dimension.
stock_code	Unique code identifying the product stock.
description	Description of the product.

dim_metadata_transactions

Column	Description
metadata_id	Primary key for the metadata transactions table.
transaction_description	Description of the transaction type or reason.
transaction_category	Category of the transaction (e.g., "sale", "return", "adjustment", "fee").

fact_sales_transactions

Column	Description
transaction_id	Primary key for the fact table.
time_id	Foreign key referencing dim_time.
location_id	Foreign key referencing dim_location.
customer_id	Foreign key referencing dim_customer.
product_id	Foreign key referencing dim_product.
metadata_id	Foreign key referencing dim_metadata_transactions.
invoice_id	Unique identifier for the invoice.
quantity	Number of units involved in the transaction (can be negative for returns).
price	Price per unit of the product (nullable; may include refunds or adjustments).

3. Packages and modules overview

3.0. Running the project

- **Requirements:** Python 3.9+ and Jupyter Notebook.
- **Use a virtual environment:** `python -m venv .venv`
- **Installation:** Run `pip install -r requirements.txt`. The requirements file includes two sets of dependencies: one specifically for the Jupyter Notebook (.ipynb) and another for the Python-based script. This separation ensures that the notebook's additional dependencies are only installed if you plan to use it, keeping the main script lightweight and efficient.
- **Create and fill a file called '.env' in the root directory of this project. Fill it with the following environment variables:**

```
MSSQL_WAREHOUSE_URL="mssql+pyodbc://<username>:<password>@<host>/<database>?  
driver=ODBC+Driver+17+for+SQL+Server&trusted_connection=yes"
```

- **Run the main script:** `python solution.py` on your terminal, from the root project directory.

3.1. Main Script

- **draw.ipynb**: Development notebook. Contains all the exploratory data analysis (EDA) and data cleaning steps. It also includes the initial data quality check and the first assumptions and abnormalities identified.
- **solution.py**: Example of a dedicated pipeline, including all modules. Possible Lineage tracker, with space to implement a run back from stopped stage.
- **.env**: Requirement to DW experience.

3.2. infra

- models
 - Dimensional and fact models.
 - **dim.py** - all dimensional to our DW models.
 - **fact.py** - all fact to our DW models.
 - **facts_integrity.py** - all base validators to our fact in the DW.
 - **dims_integrity.py** - all base validators to our dims in the DW.
- pipeline
 - Pipeline specif codes
 - **pipeline_metadata.py** - References to metadata process. Like Mapping, etc.
 - **NORMALIZE_LOCATION_MAP** - a dict containing the mapping of normalized location names.
 - **CLOUD_LOST_PRODUCTS_WORDS** - a list of words that indicate lost products.
 - **STAGE_III_COLUMNS** - renamed columns to be used in the pipeline.
 - **validation_models** - mapper containing Pydantic models to validate the data.
 - **models_map** - mapper containing sqlalchemy models to validate the data.
 - **pipeline_lineage.py** - It stores stages related to the pipeline.
 - **get_csv_df** - Reads CSV files into pandas DataFrame format.
 - **PipelineTransformer** - BR Contains every stage and their transformations, as well as a saving method.
 - **pipeline_transformers.py** - Business rules (BR) and general transformations (GR) to be used on the pipeline.
 - **sanitize_column_data** - BR related to fill null data and format types.
 - **sanitize_text** - BR related to sanitize text data. It will remove special characters, and replace accented characters with their unaccented counterparts.
 - **BaseTableGenerator** - BR related to generate base tables.
 - **DimTimeGenerator** - BR related to generate the time dimension.
 - **DimLocationGenerator** - BR related to generate the location dimension.
 - **DimCustomerGenerator** - BR related to generate the customer dimension.
 - **DimProductGenerator** - BR related to generate the product dimension.
 - **DimMetadataTransactionsGenerator** - BR related to generate the metadata transactions dimension.
 - **FactSalesTransactionsGenerator** - BR related to generate the sales transactions fact table.
 - **generate_warehouse_sales_tables** - GR related to generate the warehouse tables.
 - **validate_warehouse_sales_data** - BR related to validate the warehouse tables.
 - **validate_data_integrity** - BR related to validate the data integrity.

- handlers
 - General handlers for database related and data processing.
 - `mysql_handler.py` - MSSQL connection handler. It will be used to return the connection engine to be orchestrated by sqlalchemy/alembic/direct-queries.

3.3. ingestion

- Ingested data as 7z.
- Extracted data as csv/xls.

3.4. utils

- `_references.py` - base code references, like logging, etc.
 - `get_current_utc_time` - datetime now in UTC time.
 - `create_logger` - created supposed to use as unique logger for the application.
- `file_handlers.py` - utilities to write/read/process/format files.
 - `extract_7z` - It will extract the 7z file to a folder. This one is our "imaginary API".

3.5. assets

- Images and other assets used on README.md and other documentation. It doesn't make part of the codebase.

4. General Code Structure

- Load the data.
- Understand data scenario.
- Avoided generic changes to the data on the beginning.
- Created specific changes to the data per column.
 - Revisit different columns while exploring possible business rules.
- Transformate specific columns.
- Quick overview of transformations.
- Check data as a whole, after specific transformation.
- Apply generic changes on data.
- Validate transformations.
- Generates warehouse transformations.
- Validate warehouse transformations.
- Save data.

Code can include intermediate passes like "saving stages".

Aggregations & Reporting

I'm considering that the data is accurate enough (based on assumed prior reviews), and these represent real insights.

Aggregations.sql

1. Insight - (AVTQ): Absolute Value of Transaction Quantities and Revenues across different categories.

Magnitude of operational activity within each transaction category, without distinguishing between income and expense.

```
WITH category_totals AS (  
    SELECT  
        m.transaction_category,  
        m.transaction_description,  
        SUM(ABS(f.quantity)) AS total_sales_quantity,  
        SUM(ABS(f.quantity * f.price)) AS total_sales_revenue  
    FROM master.sales_warehousing.fact_sales_transactions AS f  
    INNER JOIN master.sales_warehousing.dim_metadata_transactions AS m  
        ON f.metadata_id = m.metadata_id  
    GROUP BY  
        m.transaction_category,  
        m.transaction_description  
)  
grand_total AS (  
    SELECT  
        SUM(ct.total_sales_revenue) AS grand_total_sales_revenue  
    FROM category_totals AS ct  
)  
SELECT  
    ct.transaction_category,  
    ct.transaction_description,  
    ct.total_sales_quantity,  
    ct.total_sales_revenue,  
    (ct.total_sales_revenue / gt.grand_total_sales_revenue) * 100 AS  
percentage_of_total_sales_revenue  
FROM category_totals AS ct  
CROSS JOIN grand_total AS gt  
ORDER BY  
    percentage_of_total_sales_revenue DESC;
```

alization
onstrates
Absolute
me and
ntity
saction
2T)
ysis
lts. This
ario
lights
nders like
luct
ns,
erating

hts into
we could
each
cing the
n rate by,
rstance,
, focusing
re top
reasons.
egies
it include
ering
omer
back,
erstanding
ons for
ns, and so
product
ity checks,
omer
ce
vements,
other
egies
d be
emented.

	transaction_category	transaction_description	total_sales_quantity	total_sales_revenue	percentage_of_total_sales_reve
1	sale	Sale	610605	1269941.22	65.922900
2	return	Return	215989	489517.86	25.410900
3	financial adjustment	Financial Adjustment	400	157169.30	8.158600
4	maintenance adjustment	Maintenance Adjustment	401	9628.98	0.499800
5	lost sale	Lost Sale	1056	146.34	0.007500

2. Insight - (SLICR): Sales Location Impact Comparison on Revenue

Monthly contribution of sales transactions by location, expressed as a percentage of the total absolute revenue for all sales transactions. A transaction is a valid sale when transaction_category = 'sale'.

```
WITH location_sales AS (  
  SELECT  
    l.location_name,  
    t.year,  
    t.month,  
    SUM(f.quantity) AS total_sales_quantity,  
    SUM(f.quantity * f.price) AS total_sales_revenue  
  FROM master.sales_warehousing.fact_sales_transactions AS f  
  INNER JOIN master.sales_warehousing.dim_location AS l  
    ON f.location_id = l.location_id  
  INNER JOIN master.sales_warehousing.dim_time AS t  
    ON f.time_id = t.time_id  
  INNER JOIN master.sales_warehousing.dim_metadata_transactions AS m  
    ON f.metadata_id = m.metadata_id  
  WHERE m.transaction_category = 'sale'  
  GROUP BY  
    l.location_name, t.year, t.month
```

```
),
total_revenue AS (
    SELECT
        SUM(ABS(total_sales_revenue)) AS absolute_total_revenue
    FROM location_sales
)
SELECT
    ls.location_name,
    ls.year,
    ls.month,
    ls.total_sales_quantity,
    ls.total_sales_revenue,
    ABS(ls.total_sales_revenue) AS absolute_sales_revenue,
    (ABS(ls.total_sales_revenue) / tr.absolute_total_revenue) * 100 AS
absolute_percentage_of_total
FROM location_sales AS ls
CROSS JOIN total_revenue AS tr
ORDER BY
    absolute_percentage_of_total DESC,
    ls.year DESC,
    ls.month DESC;
```

This visualization demonstrates the (SLICR): Sales Location Impact Comparison on Revenue analysis results. This scenario highlights locations with a higher percentage of sales revenue. Based on these insights, we can create targeted promotions and enhance marketing strategies in strong locations while identifying opportunities to explore untapped potential in weaker locations. A solid approach would involve analyzing the factors contributing to the performance of both strong and weak locations, followed by implementing appropriate campaigns and lead generation strategies.

	location_name	year	month	total_sales_quantity	total_sales_revenue	absolute_sales_revenue	absolute_percentage_of_total
1	United Kingdom	2010	10	60190	147035.27	147035.27	11.578100
2	United Kingdom	2010	11	71143	125453.09	125453.09	9.878600
3	United Kingdom	2010	9	50009	119630.21	119630.21	9.420100
4	United Kingdom	2010	6	46927	111493.93	111493.93	8.779400
5	United Kingdom	2010	3	51339	108108.82	108108.82	8.512800
6	United Kingdom	2009	12	46347	81325.14	81325.14	6.403800
7	United Kingdom	2010	4	33233	76377.71	76377.71	6.014200
8	United Kingdom	2010	7	32004	69954.32	69954.32	5.508400
9	United Kingdom	2010	8	38028	63387.85	63387.85	4.991400
10	United Kingdom	2010	1	26618	57497.48	57497.48	4.527500
11	United Kingdom	2010	5	48219	55379.50	55379.50	4.360700
12	United Kingdom	2010	12	19916	51586.79	51586.79	4.062100
13	United Kingdom	2010	2	26571	43960.46	43960.46	3.461600
14	Norway	2010	3	2	13916.34	13916.34	1.095800
15	Ireland	2010	10	750	9384.05	9384.05	0.738900
16	Ireland	2010	1	563	7420.23	7420.23	0.584200
17	France	2010	7	203	4352.45	4352.45	0.342700
18	Japan	2010	12	1496	3828.40	3828.40	0.301400
19	Ireland	2010	9	1107	3806.16	3806.16	0.299700

3. Insight - (CLV): Customer Lifetime Value

Understand most valuable customers and their revenues over their lifetime. Considered total revenue generated (lifetime_value) and engagement (months_active) for known customers, without filtering by location or transaction

type. Valid sales are those with transaction_category = 'sale'. Not distinguishing if they are Legal Entities or Individuals.

```
WITH customer_lifetime_value AS (
  SELECT
    c.customer_id,
    SUM(CASE WHEN m.transaction_category = 'sale' THEN f.quantity * f.price ELSE 0
END) AS lifetime_value,
    COUNT(DISTINCT CONCAT(t.year, '-', t.month)) AS months_active,
    CASE WHEN SUM(CASE WHEN m.transaction_category = 'sale' THEN 1 ELSE 0 END) > 0
THEN 1 ELSE 0 END AS valid_sale_flag
  FROM master.sales_warehousing.fact_sales_transactions AS f
  INNER JOIN master.sales_warehousing.dim_customer AS c
    ON f.customer_id = c.customer_id
  INNER JOIN master.sales_warehousing.dim_time AS t
    ON f.time_id = t.time_id
  INNER JOIN master.sales_warehousing.dim_metadata_transactions AS m
    ON f.metadata_id = m.metadata_id
  WHERE c.is_known_customer = 1
  GROUP BY c.customer_id
)
SELECT
  customer_id,
  lifetime_value,
  months_active,
  lifetime_value / NULLIF(months_active, 0) AS average_monthly_value
FROM customer_lifetime_value
ORDER BY lifetime_value DESC;
```

This visualization demonstrates the Customer Lifetime Value (CLV) analysis results. This scenario highlights the most valuable customers based on their lifetime value and engagement. By identifying these high-value customers, we can develop targeted marketing strategies to enhance customer loyalty and retention. Strategies might include personalized offers, loyalty programs, and exclusive promotions to increase customer engagement and lifetime value. Additionally, for instance, we can associate this with the first insight to understand how customers engage with the company and

	customer_id	lifetime_value	months_active	average_monthly_value
1	3cdacdba0a3e4bf931f5e577ee4f94f7	56807.36	13	4369.796923
2	80789d636d68ec8ac889de80365bbd57	52383.60	13	4029.507692
3	21bfef81b08fa7988c78190cc68c241c	19543.18	13	1503.321538
4	1671f1468c94ff1c45e249f5fd322ed5	18415.36	13	1416.566153
5	d4c074c60c57087c95fd0a17f986210c	17918.64	11	1628.967272
6	8a7c03958cbbb5d374f4be72690ca7e	16403.77	9	1822.641111
7	42ae58924107d3be8c6b7b7e3cfd4164	16122.29	13	1240.176153
8	1a62026e6b035b51682672932876a119	14689.84	4	3672.460000
9	7842858ddd53cb0a24dc8c9fea4f92b	14550.74	12	1212.561666
10	dfc0a2d63b0d7a1ce1cd07ffe3a3aea7	14550.13	12	1212.510833
11	01c8a64a2b3c66c05c2dbf9df27510eb	14546.26	13	1118.943076
12	0783683c446cf52f9df7d90d92bf5239	14215.71	12	1184.642500
13	ca87e9d08238e6d61791dc55931bb500	13916.34	1	13916.340000
14	4218470a524ef1991202bb63abee5d72	13285.25	8	1660.656250
15	52c031e023c8a03f30f57246f9c3d4f9	10953.50	1	10953.500000
16	1221132d8390ea66832cf2eabd8eb668	10109.40	5	2021.880000
17	74367f78d716836cc099eeb6497f8703	8950.84	12	745.903333
18	f916b2ed383e62ec91b915de8ba77e0b	8897.00	13	684.384615
19	2779d140faa92aa0ad8df4230aca4590	7715.04	13	593.464615

its offered products based on their location, helping to retain new leads and maintain existing ones, enhancing the overall customer experience.

4. Data-Driven Deep

Our approaches are guided by business rules (market, client, company, etc) and the data available to us. These factors shape the reports we generate and follow, ensuring alignment with our strategic goals.

- The first insight offers a broad overview of events within the specified date range. From this starting point, we could analyze monthly trends to uncover the seasonality of key pain points. By consistently delivering excellent services to our clients, we gain deeper insights into how to remain competitive and innovative.
- The second insight delves into location-specific impacts on revenue, providing a detailed perspective. This approach helps identify the most profitable locations to prioritize for investment, as well as underperforming areas that require strategic intervention. By continuously monitoring and taking timely action, we can prevent strong locations from weakening and improve weaker locations' performance.
- The third insight focuses on customer lifetime value (CLV), a critical metric for understanding customer engagement and loyalty. By identifying high-value customers and their engagement levels, we can leverage other insights to develop targeted strategies. I can think in a broad range of strategies that could include directly engaging with these clients, implementing new marketing approaches, or introducing new products to further engage them, while also strengthening the brand to attract new customers.

In data we trust

Tunning Scenario

This will not be performed, but I'm assuming that I would use these aggregations to create reports for the business team. Here would be my approach for also tuning this scenario.

This is a first basic approach. I'm assuming that there aren't logs yet (which could help identify most used columns, for instance).

- Tuning approach, client side:
 - Query leveraging predicate filters.
 - Understand the Logical plan by the optimizer:
 - Identify bottlenecks in the execution plan by the optimizer.
 - How could I induce the optimizer to better optimize the query plan?
 - Understanding query optimal scenario tuning.
 - Understanding query hints.

Now, after having the execution plan and a proposed scenario, I can proceed to basic server-side tuning (assuming we are on a hybrid ecosystem).

- Indexes: They need special considerations—the trade-off between maintenance and performance. Understand base columns to be indexed

- Understand base columns to be indexed
- Statistics: They are crucial to the optimizer. There is no logical in having indexes without updated statistics.
 - Governance to set frequent statistics updates.
- Partitioning, if necessary.
- Cost Threshold, parallelism, and other high-level server configurations.

We would need approaches like governance, to ensure optimal database performance. Consumption governance (help with querie base good practices). Among others.

This sales Retail refers to the Attam Store.

Matheus Alves @ 2024